


Peak device memory scaling



—●— PyTorch baseline -■- Phase 3 (Triton + dW) -▲- Phase 4 (in-kernel Philox) -◆- Phase 5 (fused reduction)